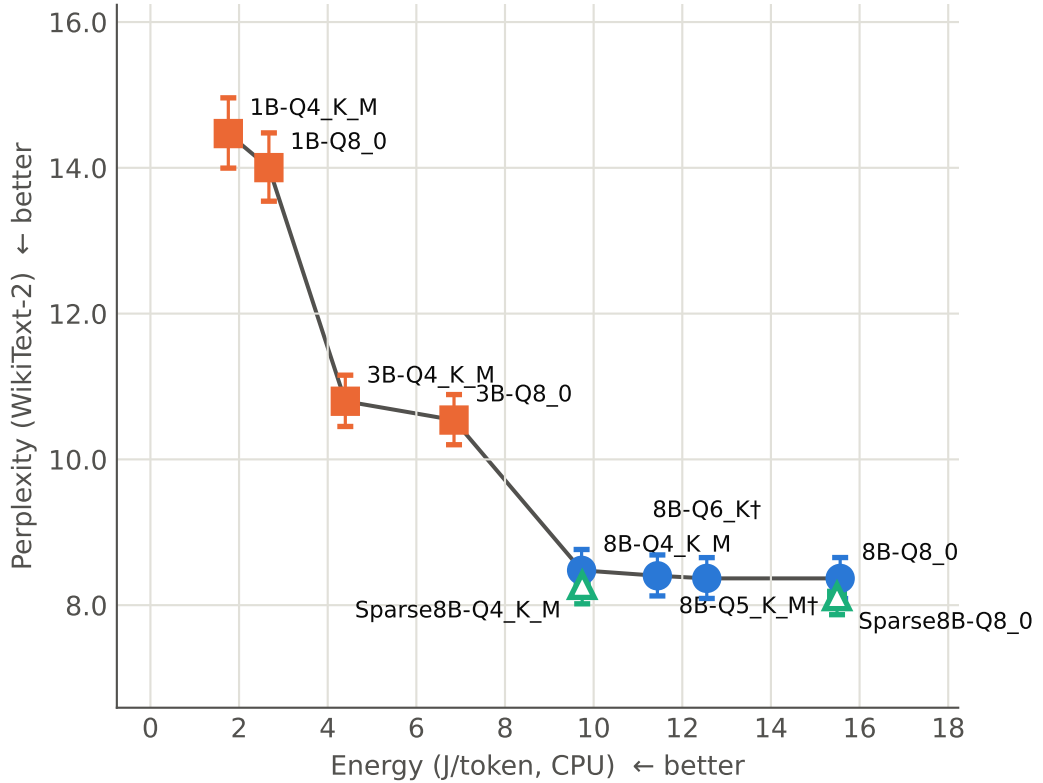

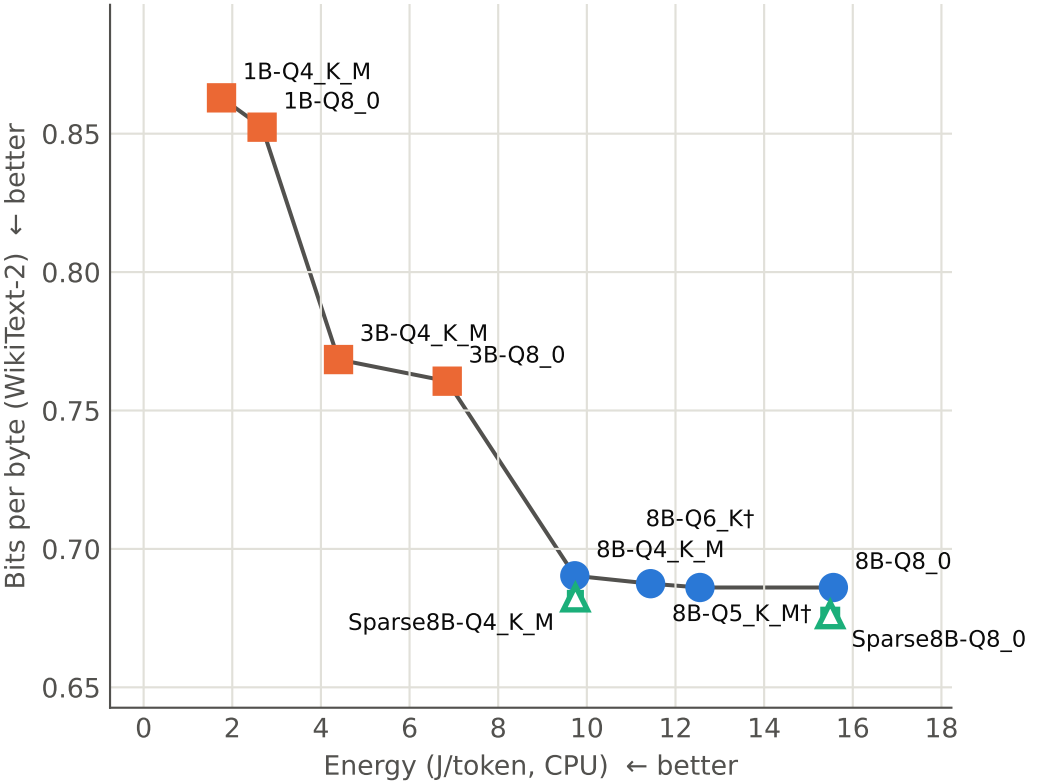

quantization · pruning + distillation · 2:4 sparse pruning + distillation · base model · Pareto frontier (instruct models)

Horizontal error bar: between-session ±1 SD (n = 3). Vertical error bar: chunk-level 95 % CI of the perplexity (564 chunks, the whole WikiText-2 test set).
The Pareto frontier is drawn over instruct models only; the base model is hollow. † 8B-Q5_K_M and 8B-Q6_K were measured in the same night (night_01, 10 rounds);
their J/token is scaled to the campaign through the anchor models (8B-Q4_K_M, 8B-Q8_0).