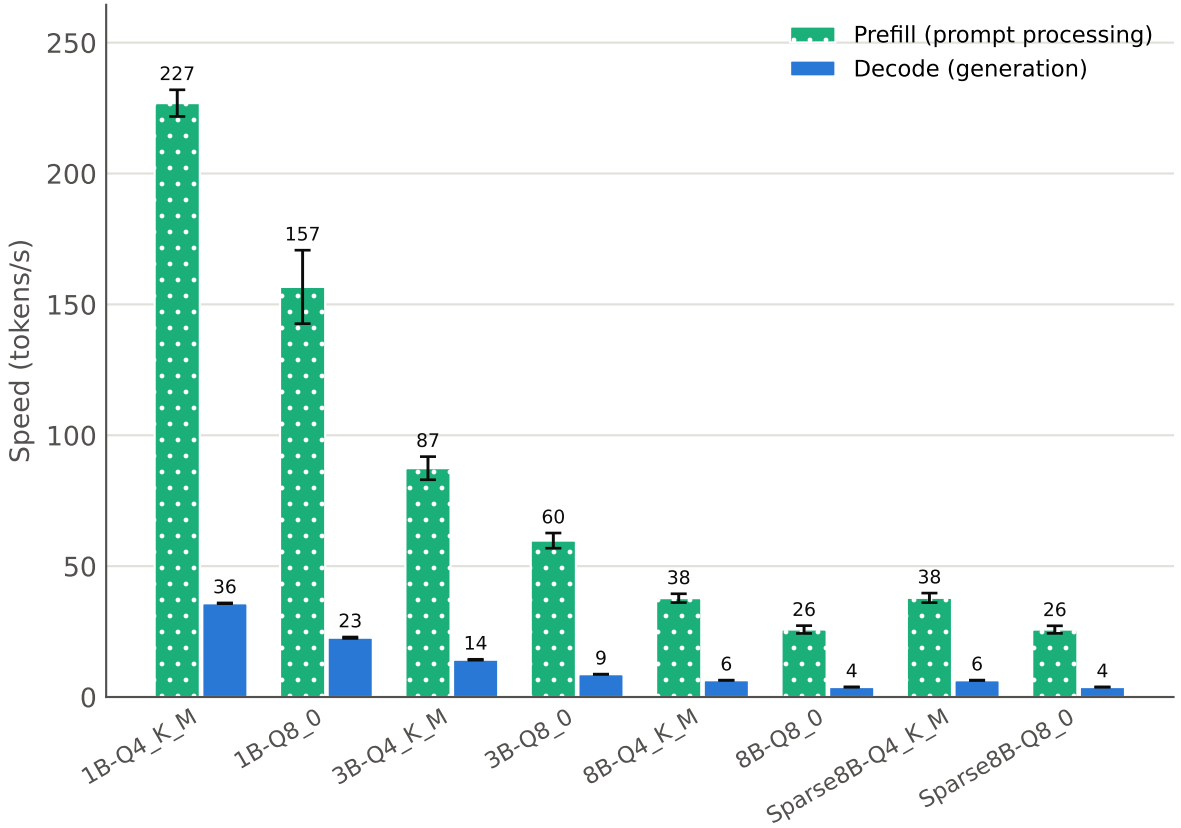

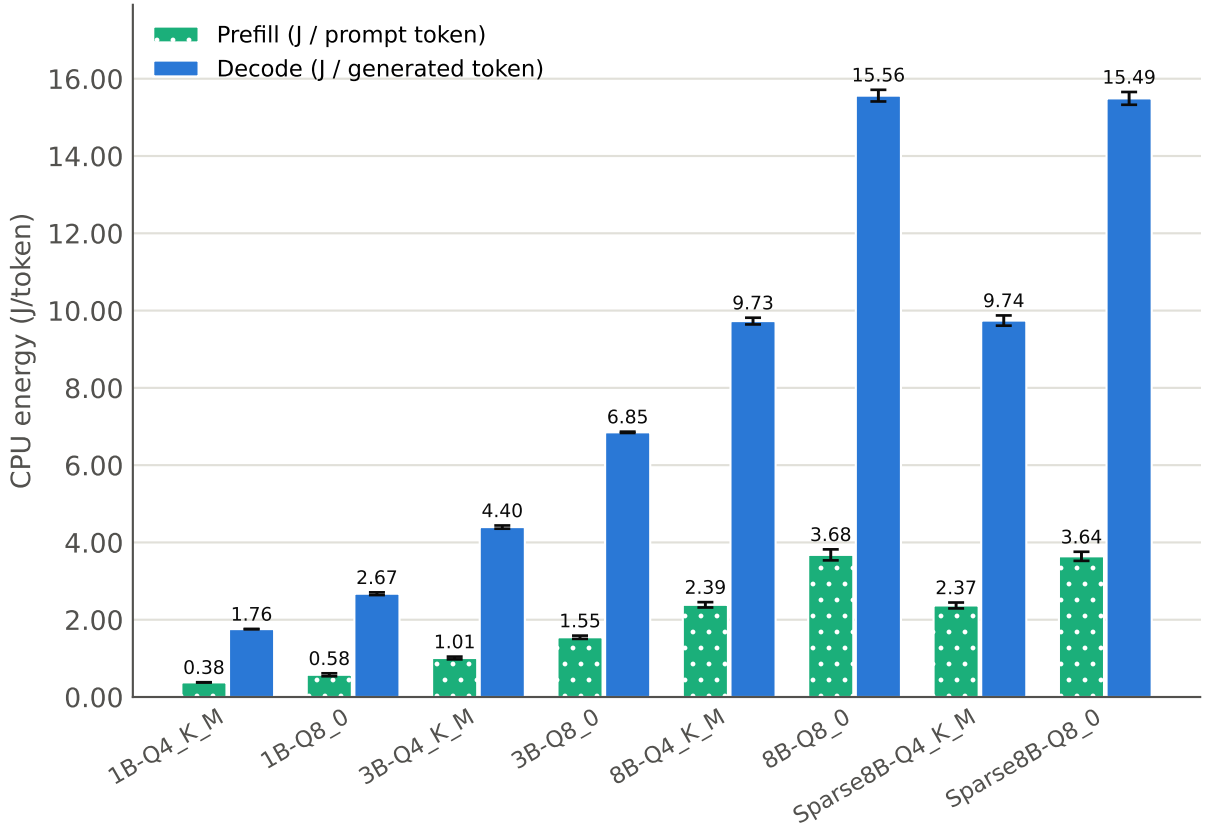

Bar: mean of three sessions; error bar: between-session ±1 SD (n = 3). Prefill is measured on 11–18-token prompts (one forward pass, compute-bound); decode is the 100-token generation. Right panel: CPU energy of the prefill window per prompt token next to the decode J/token. 8 threads.