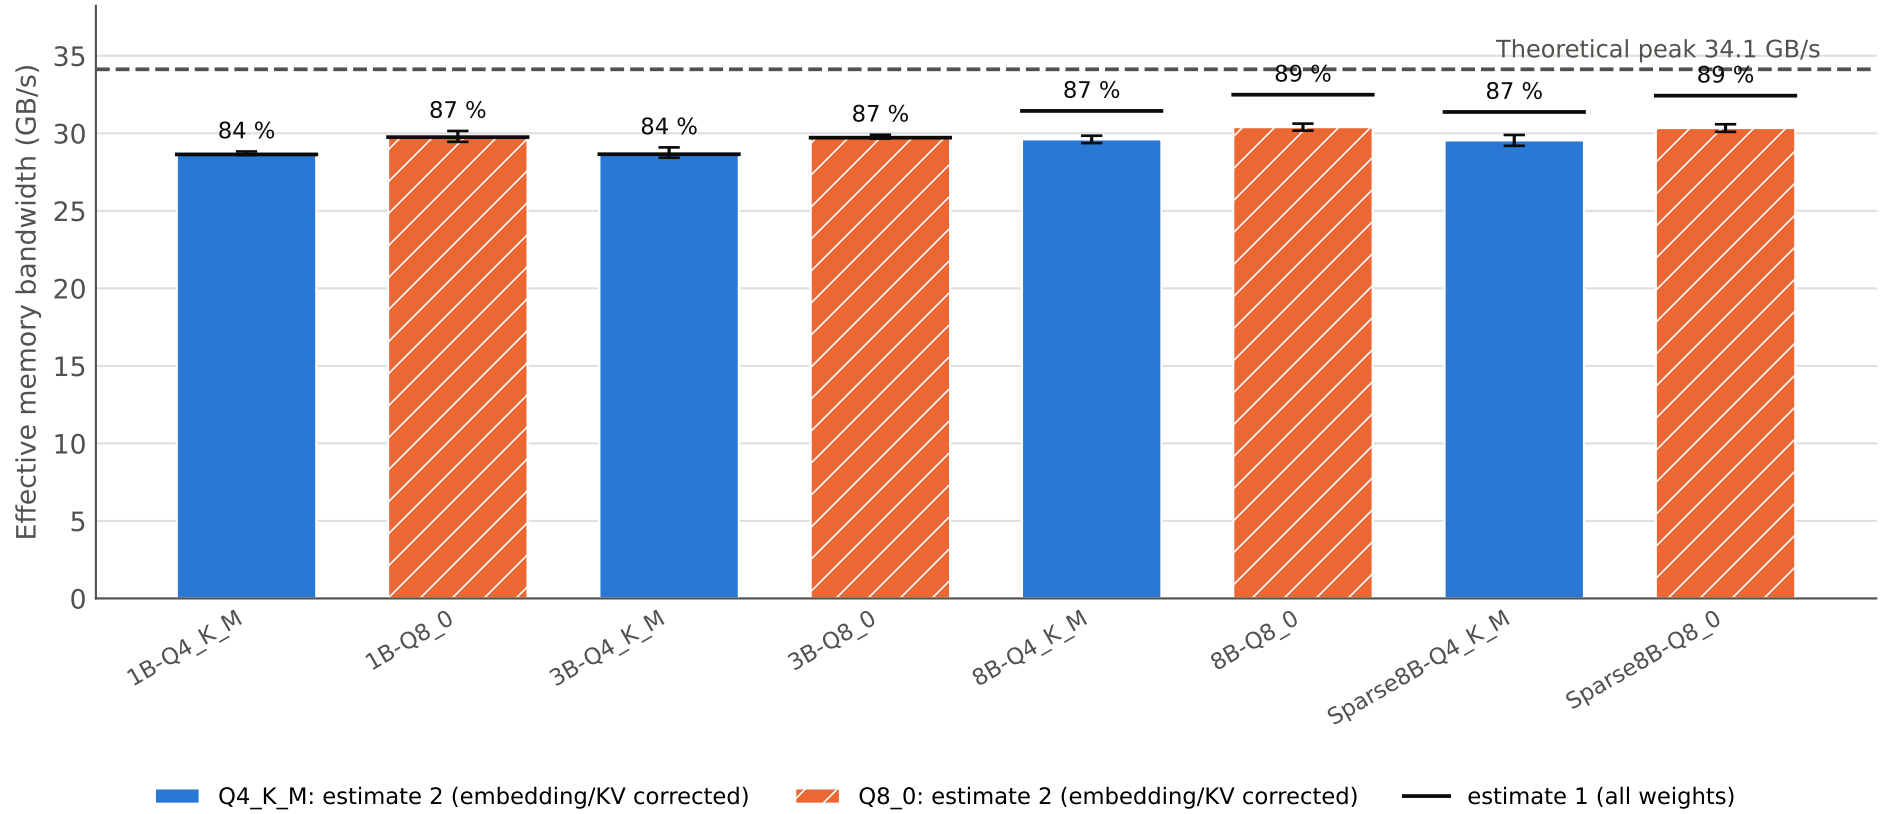

Bar: corrected estimate (weights − row-wise embedding table + KV cache) × decode tokens/s (mean of three sessions, 8 threads); error bar: between-session ±1 SD. Tick: all weights read every token (estimate 1). Percentages are the share of the theoretical peak.